

# Intrusion Detection Analysis

October 7, 2018

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

## 0.1 1. Introduce the Data

### 0.1.1 Import the dataset

dataset.txt is a tab delimited file that has around 5 lac rows and 24 features. We read the data into an object called df and then assign the labels of the features to their respective columns in the dataframe df.

```
In [2]: df = pd.read_csv('dataset.txt', sep = "\t", header = None) #use nrow attribute for limit
column_list = [
    'duration',
    'service',
    'source_bytes',
    'destination_bytes',
    'count',
    'same_srv_rate',
    'serror_rate',
    'srv_serror_rate',
    'dst_host_count',
    'dst_host_srv_count',
    'dst_host_same_src_port_rate',
    'dst_host_serror_rate',
    'dst_host_srv_serror_rate',
    'flag',
    'ids_detection',
    'malware_detection',
    'ashula_detection',
    'label',
    'source_ip_address',
    'source_port_number',
    'destination_ip_address',
    'destination_port_number',
    'start_time',
```

```

        'protocol'
    ]
    df.columns = column_list

```

```

/home/chrx/anaconda3/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2785: DtypeWarning:
  interactivity=interactivity, compiler=compiler, result=result)

```

In [3]: *#reading the first 5 rows of the data*

```
df.head()
```

```

Out[3]:
  duration service  source_bytes  destination_bytes  count  same_srv_rate  \
0  0.000393    dns           43           100          9          1.0
1  0.000400    dns           43           100         11          1.0
2  0.000432    dns           61            77         14          1.0
3  0.000442    dns           61            77         15          1.0
4  0.000452    dns           61           112         16          1.0

  serror_rate  srv_serror_rate  dst_host_count  dst_host_srv_count  ...  \
0          0.0              0.0             60                 99    ...
1          0.0              0.0             61                 99    ...
2          0.0              0.0             60                 99    ...
3          0.0              0.0             39                 99    ...
4          0.0              0.0             39                 99    ...

  ids_detection  malware_detection  ashula_detection  label  \
0              0                  0                  0     -1
1              0                  0                  0     -1
2              0                  0                  0     -1
3              0                  0                  0      1
4              0                  0                  0      1

  source_ip_address  source_port_number  \
0  fd95:ec1e:6a61:df6b:7de2:27ad:6105:3709    46071
1  fd95:ec1e:6a61:df6b:7de2:27ad:6105:3709    51986
2  fd95:ec1e:6a61:df6b:7de2:27ad:6105:3709    48133
3  fd95:ec1e:6a61:b804:7dcf:276f:0751:0ff5    58986
4  fd95:ec1e:6a61:b804:7dcf:276f:0751:0ff5    35589

  destination_ip_address  destination_port_number  \
0  fd95:ec1e:6a61:435f:7de5:27b0:7d84:3c0d      53
1  fd95:ec1e:6a61:435f:7de5:27b0:7d84:3c0d      53
2  fd95:ec1e:6a61:435f:7de5:27b0:7d84:3c0d      53
3  fd95:ec1e:6a61:435f:7de5:27b0:7d84:3c0d      53
4  fd95:ec1e:6a61:435f:7de5:27b0:7d84:3c0d      53

  start_time  protocol

```

```

0    00:00:00      udp
1    00:00:00      udp
2    00:00:00      udp
3    00:00:00      udp
4    00:00:00      udp

```

```
[5 rows x 24 columns]
```

```
In [4]: #shape of the data is the number of rows by the number of features
df.shape
```

```
Out[4]: (499999, 24)
```

### 0.1.2 Remove the target variable from the dataset

the target variable will be label

**Important** -- label = 0 if no Intrusion 1 Otherwise

although malware detection, ids\_detection, and ashula\_detection could also be taken as labels, we decided against it since they were only indicative of the detection of intrusion by a software and therefore may not be correctly labelled.

```
In [5]: df['label'] = [0 if x == 1 else 1 for x in df['label']]
```

```
In [6]: label_target = df.pop('label').values
```

```

#dropping these labels because they aren't really needed in training
df.pop('ids_detection').values
df.pop('malware_detection').values
df.pop('ashula_detection').values

df.shape

```

```
Out[6]: (499999, 20)
```

```
In [7]: #creating a variable y that contains the target column `label` for the training set
y = label_target
```

### 0.1.3 Checking String Based Features

Here we check the unique values of string based features and determine if they are necessary. If there are a lot of unique values then the feature is ignored since it might lead to overfitting of the model

```
In [8]: #count the unique values in service feature(string based)
service_value_counts = df['service'].value_counts()
print("Number of unique values = ", service_value_counts.shape[0], "\n")
print(service_value_counts)
```

Number of unique values = 12

dns	313014
other	114253
ssh	71349
sip	617
snmp	199
smtp	182
radius	157
rdp	94
http	89
smtp,ssl	33
dhcp	11
ssl	1

Name: service, dtype: int64

```
In [9]: #count the unique values in protocol feature(string based)
        protocol_value_counts = df['protocol'].value_counts()
        print("Number of unique values = ", protocol_value_counts.shape[0], "\n")
        print(protocol_value_counts)
```

Number of unique values = 3

udp	320156
tcp	177207
icmp	2636

Name: protocol, dtype: int64

```
In [10]: #count the unique values in flag feature(string based)
         protocol_value_counts = df['flag'].value_counts()
         print("Number of unique values = ", protocol_value_counts.shape[0], "\n")
         print(protocol_value_counts)
```

Number of unique values = 12

SF	378699
S0	101335
OTH	9414
REJ	5520
RSTO	2250
SHR	1522
RSTRH	659
RSTOSO	320
S1	152
RSTR	104
S2	13
SH	11

Name: flag, dtype: int64

#### 0.1.4 Unique values for each feature in the dataset

```
In [11]: df.T.apply(lambda x: x.nunique(), axis=1)
```

```
Out[11]: duration          94392
         service            12
         source_bytes       490
         destination_bytes  371
         count             101
         same_srv_rate      54
         serror_rate        80
         srv_serror_rate    101
         dst_host_count     101
         dst_host_srv_count 101
         dst_host_same_src_port_rate 40
         dst_host_serror_rate 98
         dst_host_srv_serror_rate 101
         flag               12
         source_ip_address  12815
         source_port_number 60971
         destination_ip_address 974
         destination_port_number 3148
         start_time         74672
         protocol           3
         dtype: int64
```

#### 0.1.5 Removing unnecessary features

Everything from source\_ip\_address to start\_time is of no use because these things are really random...

```
In [12]: df.pop('source_ip_address').values
         df.pop('source_port_number').values
         df.pop('destination_ip_address').values
         df.pop('destination_port_number').values
         df.pop('start_time').values
         df.head()
```

```
Out[12]:
```

	duration	service	source_bytes	destination_bytes	count	same_srv_rate	\
0	0.000393	dns	43	100	9	1.0	
1	0.000400	dns	43	100	11	1.0	
2	0.000432	dns	61	77	14	1.0	
3	0.000442	dns	61	77	15	1.0	
4	0.000452	dns	61	112	16	1.0	

```

         serror_rate  srv_serror_rate  dst_host_count  dst_host_srv_count  \
```

0	0.0	0.0	60	99
1	0.0	0.0	61	99
2	0.0	0.0	60	99
3	0.0	0.0	39	99
4	0.0	0.0	39	99

	dst_host_same_src_port_rate	dst_host_serror_rate	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	

	dst_host_srv_serror_rate	flag	protocol
0	0.0	SF	udp
1	0.0	SF	udp
2	0.0	SF	udp
3	0.0	SF	udp
4	0.0	SF	udp

### 0.1.6 Features to use

```
In [13]: list(df)
```

```
Out[13]: ['duration',
          'service',
          'source_bytes',
          'destination_bytes',
          'count',
          'same_srv_rate',
          'serror_rate',
          'srv_serror_rate',
          'dst_host_count',
          'dst_host_srv_count',
          'dst_host_same_src_port_rate',
          'dst_host_serror_rate',
          'dst_host_srv_serror_rate',
          'flag',
          'protocol']
```

### 0.1.7 What the data looks like now in terms of data type

```
In [14]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 499999 entries, 0 to 499998
Data columns (total 15 columns):
duration          499999 non-null float64
service           499999 non-null object
```

```

source_bytes          499999 non-null int64
destination_bytes     499999 non-null int64
count                 499999 non-null int64
same_srv_rate         499999 non-null float64
serror_rate           499999 non-null float64
srv_serror_rate       499999 non-null float64
dst_host_count        499999 non-null int64
dst_host_srv_count    499999 non-null int64
dst_host_same_src_port_rate 499999 non-null float64
dst_host_serror_rate  499999 non-null float64
dst_host_srv_serror_rate 499999 non-null float64
flag                  499999 non-null object
protocol              499999 non-null object
dtypes: float64(7), int64(5), object(3)
memory usage: 57.2+ MB

```

### 0.1.8 Transfrom Catergorical Data to Numerical Data

From the above result it can be seen that service, flag, and protocol are not numeric data. Since we have to supply the machine learning models with numeric data we have to somehow transform the categorical data to numeric data.

For this purpose we use a Label Encoder that encodes the unique values of a feature to a unique numeric constant(number). We do this encoding for all the rows in the data.

```

In [15]: from sklearn import preprocessing

In [16]: #get the unique values for the following categorical data
categorical_data = ['service', 'flag', 'protocol']
unique_flag_data = df['flag'].unique()
unique_service_data = df['service'].unique()
unique_protocol_data = df['protocol'].unique()

```

**Encoder for feature : Flag** This shows an example of how categorical data like flag may be encoded and then decoded from categorical to numeric and then from numeric to categorical respectively.

```

In [17]: le_flag = preprocessing.LabelEncoder()
         #Fit the label encoder to unique values
         le_flag.fit(unique_flag_data)

         #Fit the label data to some example data
         example_flag_data = list(df.head()['flag'])
         #Fit the label encoder and return encoded labels
         encoded_flag_data = le_flag.transform(example_flag_data)

         #Transform labels back to original encoding
         decoded_flag_data = list(le_flag.inverse_transform(encoded_flag_data))

```

```

print(example_flag_data)
print(encoded_flag_data)
print(decoded_flag_data)
#Ignore any warnings

```

```

['SF', 'SF', 'SF', 'SF', 'SF']
[9 9 9 9 9]
['SF', 'SF', 'SF', 'SF', 'SF']

```

/home/chrax/anaconda3/lib/python3.6/site-packages/sklearn/preprocessing/label.py:151: DeprecationWarning:
if diff:

### Encoder for feature : service and protocol

```

In [18]: le_service = preprocessing.LabelEncoder()
         le_service.fit(unique_service_data)

         le_protocol = preprocessing.LabelEncoder()
         le_protocol.fit(unique_protocol_data)

```

Out[18]: LabelEncoder()

### Encode the categorical features for all rows in the data

```

In [19]: df['flag'] = le_flag.transform(df['flag'])
         df['service'] = le_service.transform(df['service'])
         df['protocol'] = le_protocol.transform(df['protocol'])
         df.head()

```

```

Out[19]:
  duration  service  source_bytes  destination_bytes  count  same_srv_rate  \
0  0.000393         1           43                100        9           1.0
1  0.000400         1           43                100       11           1.0
2  0.000432         1           61                 77       14           1.0
3  0.000442         1           61                 77       15           1.0
4  0.000452         1           61                112       16           1.0

  serror_rate  srv_serror_rate  dst_host_count  dst_host_srv_count  \
0           0.0             0.0             60                 99
1           0.0             0.0             61                 99
2           0.0             0.0             60                 99
3           0.0             0.0             39                 99
4           0.0             0.0             39                 99

  dst_host_same_src_port_rate  dst_host_serror_rate  \
0                        0.0                0.0
1                        0.0                0.0

```



2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

	dst_host_srv_serror_rate	flag	protocol
0	0.0	9	2
1	0.0	9	2
2	0.0	9	2
3	0.0	9	2
4	0.0	9	2

### 0.1.9 PCA

PCA stands for Principal Component Analysis. This algorithm is used for dimensionality reduction. The algorithm is supplied with the number of dimension to output and then the PCA algorithm automatically calculates the new dimensions from the old dimensions. New dimensions are really a linear combination of the old dimensions.

```
In [20]: from sklearn.decomposition import PCA
```

```
pca = PCA(n_components = 2)
X_pca = pd.DataFrame(pca.fit_transform(df))
```

```
In [21]: X_pca.head()
```

```
Out[21]:
```

	0	1
0	-241.430381	-83.353206
1	-241.431096	-83.369025
2	-264.278958	-65.191378
3	-264.274311	-65.113415
4	-229.275807	-65.414867

## 0.2 Model Building

We have used 3 different types of machine learning models which are:

- KMeans
- Logistic Regression
- Random Forests

Here KMeans is a Unsupervised Learning model and the other 2 are Supervised Learning Models.

### 0.2.1 Kmeans

Here we use 2 clusters because we want the data to cluster into 2 clusters: Intrusion or Not Intrusion

```

In [22]: from sklearn.cluster import KMeans

         #converting to numpy array : format needed by sklearn
         X_pca_np = np.array(X_pca).astype(float)

         #creating a model
         kmeans = KMeans(n_clusters = 2, random_state = 0)

         #fitting the data to the model
         kmeans.fit(X_pca_np)

         #plotting the points and the cluster centroids
         plt.scatter(X_pca_np[:,0],X_pca_np[:,1], c = kmeans.labels_, cmap = 'rainbow')
         plt.scatter(kmeans.cluster_centers_[0],kmeans.cluster_centers_[1], color='black')


         #frequency of labels
         predicted = kmeans.predict(X_pca_np)
         unique, counts = np.unique(predicted, return_counts=True)
         print(np.asarray((unique, counts)).T)
         unique, counts = np.unique(y, return_counts=True)
         print(np.asarray((unique, counts)).T)

```

```

[[ 0 499995]
 [ 1      4]]
[[ 0 129616]
 [ 1 370383]]

```



```

In [23]: #calculating the percentage of correct clusterings
from sklearn.metrics import accuracy_score
print("accuracy Score : ",accuracy_score(y, kmeans.predict(X_pca_np)))

#printing confusion matrix
from sklearn.metrics import confusion_matrix
print("Confusion Matrix\n",confusion_matrix(y, kmeans.predict(X_pca_np)))

#plotting confusion matrix
plt.imshow(confusion_matrix(y, kmeans.predict(X_pca_np)),
           cmap='Blues', interpolation='nearest')
plt.colorbar()
plt.grid(False)
plt.ylabel('true')
plt.xlabel('predicted');

```


accuracy Score : 0.25924051848103696

Confusion Matrix

```

[[129616    0]
 [370379    4]]

```



**Lets try with a change in hyperparameters to kmeans**

```

In [24]: #fitting data to modded kmeans model
kmeans_modded = KMeans(n_clusters = 2, random_state = 0, max_iter = 100, algorithm = 'a

```

```

kmeans_modded.fit(X_pca_np)

#calculating the percentage of correct labels
print("accuracy Score : ",accuracy_score(y, kmeans_modded.predict(X_pca_np)))
#printing confusion matrix
print("Confusion Matrix\n",confusion_matrix(y, kmeans_modded.predict(X_pca_np)))

#plotting confusion matrix
plt.imshow(confusion_matrix(y, kmeans_modded.predict(X_pca_np)),
           cmap='Blues', interpolation='nearest')
plt.colorbar()
plt.grid(False)
plt.ylabel('true')
plt.xlabel('predicted');

```


accuracy Score : 0.25924051848103696

Confusion Matrix

```

[[129616    0]
 [370379    4]]

```



**Lets try to scale the data and then run kmeans**

In [25]: scaler = preprocessing.MinMaxScaler()

```

#scaling the data
X_pca_np_scaled = scaler.fit_transform(X_pca_np)

#fitting the data to modded kmeans
kmeans_modded.fit(X_pca_np_scaled)

#caclulating the percentage of correct labels
print("accuracy Score : ",accuracy_score(y, kmeans_modded.predict(X_pca_np_scaled)))
#printing confusion matrix
print("Confusion Matrix\n",confusion_matrix(y, kmeans_modded.predict(X_pca_np_scaled)))

```

accuracy Score : 0.25924051848103696

Confusion Matrix

```

[[129616    0]
 [370379    4]]

```

## 0.2.2 Regression

**Splitting the dataset** We split the data into Training and Test set based on the parameter `train_size`.

Example : if `train_size = 0.70` then Training Set contains 80% of the data and Testing Set other 20%

```

In [26]: # Use train_test_split in sklearn.cross_validation to split data into train and test set
from sklearn.cross_validation import train_test_split

```

```

X_train, X_test, y_train, y_test = train_test_split(df, y, train_size=0.70, random_stat

```

```

/home/chr/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarn
"This module will be removed in 0.20.", DeprecationWarning)

```

```

In [27]: # Function to build model and find model performance
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score

def find_model_perf(X_train, y_train, X_test, y_test, func):
    model = func()
    model.fit(X_train, y_train)
    y_hat = [x[1] for x in model.predict_proba(X_test)]
    auc = roc_auc_score(y_test, y_hat)

    return auc

```

```

In [28]: # Find performance of model using preprocessed data
auc_processed = find_model_perf(X_train, y_train, X_test, y_test, LogisticRegression)
print(auc_processed)

```

0.9274834718737871

This shows that Regression can correctly predict 98.5% of data

### 0.2.3 Random Forests

```
In [29]: from sklearn.ensemble import RandomForestClassifier
         randomForest = RandomForestClassifier(n_estimators=100, random_state=0)

In [30]: randomForest.fit(X_train, y_train)

Out[30]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)


In [31]: #calculating the percentage of correct labels
         print("accuracy Score : ",accuracy_score(y_test, randomForest.predict(X_test)))
         #printing confusion matrix
         print("Confusion Matrix\n",confusion_matrix(y_test, randomForest.predict(X_test)))

         #plotting confusion matrix
         plt.imshow(confusion_matrix(y_test, randomForest.predict(X_test)),
                     cmap='Blues', interpolation='nearest')
         plt.colorbar()
         plt.grid(False)
         plt.ylabel('true')
         plt.xlabel('predicted');
```

accuracy Score : 0.9901333333333333

Confusion Matrix

```
[[ 38059   874]
 [   606 110461]]
```



#### 0.2.4 Conclusion

Random Forests is the best available predictor for intrusion detection with Accuracy Score of 99.01%